

On this page

Monitoring Guide

Case Manager uses [LetMeKnow](#) library to export different application metrics.

The metrics are available in the following HTTP endpoints:

URI	Description
<code>/api/metrics</code>	The metrics in the dropwizard format.
<code>/api/metrics/rich</code>	The metrics in the LetMeKnow format.
<code>/api/metrics/prometheus</code>	The metrics in the Prometheus format.
<code>/api/healthchecks</code>	CM healthchecks in CSV format.

Health Checks

CM is exporting the following health checks to check if the following dependencies are working properly:

- Database
- Messaging Server
- Pulse
- Tokenizer
- Event Ingestion

The Health Checks catch all the exceptions thrown and log them. If any of the healthchecks fail, please refer to the CM logs for more details on the error.

Database Health Check

The database health check relies on the shared CM's connection pool to execute the statement `SELECT 1`. If the connection is successfully retrieved and the statement finishes, the health check succeeds, if any error occurs in any of the steps the check fails. A timeout of 5 seconds was given when retrieving a connection to avoid hanging the health check response.

Messaging Health Check

The Messaging Server Health Check is achieved by trying to initialize a new connection with the messaging server.

Tokenizer Health Check

The tokenizer health checks validates if there is a tokenizer instance reachable and if the authorization configured is valid.

Pulse Health Check

The Pulse Health Check when triggered tries to connect to Pulse and tries to login, if it succeeds the check will also succeed.

Event Ingestion Health Check

The Event Health Check checks if the current Case Manager node is taking events from the event RabbitMQ, or if it is blocking this action.

Metrics

CM already exports the default metrics provided by [LetMeKnow](#). On top of these metrics, CM also exports metrics related to the following components:

Connection Pool

Metric	Description	Tags	Type
<code>cm.db.connection.pool.size</code>	The current size of the connection pool (If pool is lazy the value can grow over time)	<code>appType=CASE_MANAGER</code> <code>metricType=APPLICATION</code>	Gauge
<code>cm.db.connection.pool.used</code>	The number of used database connections.	<code>appType=CASE_MANAGER</code> <code>metricType=APPLICATION</code>	Gauge
<code>cm.db.connection.pool.free</code>	The number of free database connections.	<code>appType=CASE_MANAGER</code> <code>metricType=APPLICATION</code>	Gauge
<code>cm.db.connection.pool.waiting</code>	The number of threads waiting for a database connection	<code>appType=CASE_MANAGER</code> <code>metricType=APPLICATION</code>	Gauge
<code>cm.db.connection.pool.usage</code>	The current usage of the database connection pool (used/size in percent)	<code>appType=CASE_MANAGER</code> <code>metricType=APPLICATION</code>	PercentGauge

Event Ingestion

The latency metric has 3 dimensions specified by the tag "eventType" each dimension refers to the type of the event the metric is measuring.

Metric	Description	Tags	Type
<code>cm.event.processing.deadletter.count</code>	The total number of dead letters since CM's startup.	<code>appType=CASE_MANAGER</code> <code>metricType=APPLICATION</code>	Counter
<code>cm.event.processing.latency</code>	The event processing latency and rate aggregated over a 10 minute sliding window.	<code>appType=CASE_MANAGER</code> <code>metricType=APPLICATION</code> <code>eventType={NEW_EVENT, STATUS_CHANGE, NEW_EXPLANATION}</code>	Timer
<code>cm.event.processing.queue.used.count</code>	The number of events waiting to be processed.	<code>appType=CASE_MANAGER</code> <code>metricType=APPLICATION</code>	Gauge
<code>cm.event.processing.queue.ratio</code>	The percentage of the event processing queue being used.	<code>appType=CASE_MANAGER</code> <code>metricType=APPLICATION</code>	Gauge

CM API

The metrics below can have several dimensions defined by the tag "path" which indicates the URI of the API the metric is measuring.

Metric	Description	Tags	Type
<code>cm.request.api</code>	The latency of the endpoint described in the path tag, aggregated in a 10 minutes bucket.	<code>appType=CASE_MANAGER</code> <code>metricType=APPLICATION</code>	Timer

Metric	Description	Tags	Type
		path={ The API path the metrics refers to}	
cm.request.api.error	The rate of errors returned by the endpoint described in the path tag.	appType=CASE_MANAGER metricType=APPLICATION path={ The API path the metrics refers to}	Meter

Tokenizer / Encryption

The metrics below can have two dimensions defined by the tag "operation" which indicates the operation the metric is measuring.

Metric	Description	Tags	Type
cm.request.tokenizer	The latency of the Tokenizer operation described, aggregated in a 10 minutes bucket.	appType=CASE_MANAGER metricType=APPLICATION operation={ The operation the metrics refers to (getToken, getPlainText, getBulkToken, getBulkPlainText, encryptData)}	Timer
cm.request.tokenizer.error	The number of errors returned by the Tokenizer operation described.	appType=CASE_MANAGER metricType=APPLICATION operation={ The operation the metrics refers to (getToken, getPlainText, getBulkToken, getBulkPlainText, encryptData)}	Gauge
cm.request.tokenizer.error.percent	The percentage of errors returned by the Tokenizer operation described.	appType=CASE_MANAGER metricType=APPLICATION operation={ The operation the metrics refers to (getToken, getPlainText, getBulkToken, getBulkPlainText, encryptData)}	PercentGauge

Pulse Requests

The metrics below can have several dimensions defined by the tag "path" which indicates the URI of the API the metric is measuring.

Metric	Description	Tags	Type
cm.request.pulse.api	The latency of the Pulse endpoint described in the path tag, aggregated in a 10 minutes bucket.	appType=CASE_MANAGER metricType=APPLICATION path={ The API path the metrics refers to}	Timer
cm.request.pulse.api.error	The rate of errors returned by the Pulse endpoint described in the path tag.	appType=CASE_MANAGER metricType=APPLICATION path={ The API path the metrics refers to}	Meter

Automation Rules^â

The metrics below can have several dimensions defined by the tag "ruleName" which indicates the Automation Rule the metric is measuring.

Metric	Description	Tags	Type
cm.automation	The latency of the automation execution.	appType=CASE_MANAGER metricType=APPLICATION ruleName={ The rule name the metrics refers to}	Timer
cm.automation.error	The rate of errors returned by the automation execution.	appType=CASE_MANAGER metricType=APPLICATION ruleName={ The rule name the metrics refers to}	Meter

Login^â

cm.request.login	The number of login requests.	appType=CASE_MANAGER metricType=APPLICATION	Gauge
cm.request.login.2fa	The number of 2fa challenges requested during login requests.	appType=CASE_MANAGER metricType=APPLICATION	Gauge
cm.request.login.error	The number of login requests with errors.	appType=CASE_MANAGER metricType=APPLICATION	Gauge
cm.request.login.error.percent	The current percentage of login requests with errors.	appType=CASE_MANAGER metricType=APPLICATION	PercentGauge

CM's Worker Threads - Event Ingestion^â

cm.executor.pool.activeThreadsCount	The total number of active threads in the executor.	appType=CASE_MANAGER metricType=APPLICATION	Gauge
cm.executor.pool.taskQueueEventsCount	The total number of tasks submitted to the executor and pending for execution.	appType=CASE_MANAGER metricType=APPLICATION	Gauge
cm.executor.pool.maxPoolSize	The current size of the thread pool executor.	appType=CASE_MANAGER metricType=APPLICATION	Gauge
cm.executor.pool.threadsUsageRatio	The current usage of the thread pool executor.	appType=CASE_MANAGER metricType=APPLICATION	PercentGauge

CM's Worker Threads - Automation Rules Processing

<code>cm.executor.automations.pool.activeThreadsCount</code>	The total number of active threads in the executor.	appType=CASE_MANAGER metricType=APPLICATION	Gauge
<code>cm.executor.automations.pool.taskQueueEventsCount</code>	The total number of tasks submitted to the executor and pending for execution.	appType=CASE_MANAGER metricType=APPLICATION	Gauge
<code>cm.executor.automations.pool.maxPoolSize</code>	The current size of the thread pool executor.	appType=CASE_MANAGER metricType=APPLICATION	Gauge
<code>cm.executor.automations.pool.threadsUsageRatio</code>	The current usage of the thread pool executor.	appType=CASE_MANAGER metricType=APPLICATION	PercentGauge

CM's Batches

Case Manager uses several batches to persist data in the database. The following metrics measure the batch's queue occupancy percentage and the flush latency.

<code>cm.event.processing.queue.occupancy.[batch-name]</code>	The occupancy percentage of batch's queue. This queue is used as a buffer before sending events to the batch.	appType=CASE_MANAGER metricType=APPLICATION	PercentGauge
<code>cm.event.processing.batch.flush.latency.[batch-name]</code>	The latency to flush a batch to the database in entries/s.	appType=CASE_MANAGER metricType=APPLICATION	Timer

Dead Letter Metrics

Learn more on the dead letter metrics in the [Dead Letters System](#) page.

Async Extensions

The following metrics can have several dimensions defined by the tag "extension" which indicates the name of the Extension the metric is measuring.

Metric	Description	Tags	Type
<code>cm.extension</code>	The latency of the async extension execution.	appType=CASE_MANAGER metricType=APPLICATION extension={ The extension name the metrics refers to }	Timer
<code>cm.extension.error</code>	The rate of errors returned by the extensions execution.	appType=CASE_MANAGER metricType=APPLICATION extension={ The extension name the metrics refers to }	Meter

Learn more on [Monitoring Extensibility](#).

Webhook Metrics[↗](#)

The following metrics can have several dimensions defined by the tag URL which indicates the webhook URL that the metric is measuring.

Metric	Description	Tags	Type
cm.webhook.total	The total of executed webhook requests.	appType=CASE_MANAGER metricType=APPLICATION url={ The webhook URL }	Counter
cm.webhook.error	The total of failed webhook requests.	appType=CASE_MANAGER metricType=APPLICATION url={ The webhook URL }	Counter