

On this page

What's Changed

Release 0.15.1 [â](#)

Allow configuration of JVM Opts per environment [â](#)

Rework

We need JVM Opts to be set on the per each environment - i.e.: we don't expect DEV environments to require as much memory as Prod environments.

Previously, the service environment variables were set in the `<component>.install` ansible tags only, and since tasks with this tag are used to generate environment agnostic base images, service environment variables were not overridden per environment.

Fix "Load Historical Data by ETL" task tags [â](#)

Bug Fix

The task that actually launched the ETL tool was mistakenly tagged as `etl.install`. This actually meant ETL started when building the base images through Packer, which naturally fails.

We've changed the tag to `etl.configure`, which is actually the tag that runs when an instance starts - allowing for environment specific configurations and to actually run our services.

Add Project ETL configurations artifact to documentation [â](#)

Docs

Adds Project ETL Configurations to our [artifacts](#) page. Even though this artifact already existed it wasn't clearly listed as one of the installation dependencies.

Release 0.15.0 [â](#)

Allow Ansible Role to call RabbitMQ federation API without SSH [â](#)

Feature

There was an outstanding item from our Blue-Green proposal that, while it wasn't a blocker, it would mean real-time metrics could become inaccurate in that operation.

There was already a playbook to setup federation for the blue-green but the playbook assumed an SSH connection to a RabbitMQ Server host and running `rabbitmqctl` there.

With this new change `rabbitmqctl` can be used from a remote instance without SSH access. [Blue-Green Documentation](#) has been updated accordingly.

Add guardrails to avoid submission of events too far in the past or future [â](#)

Feature

In order to avoid future occurrences of the current production issue, we're deploying guardrails so that our API rejects events too far in the past or future - in our case, 30 minutes.

We want to reject events too far in the past since those would effectively induce errors in the current metrics (e.g.: an event sent with timestamp of "2 days before" has to be rejected so that it doesn't increase the count of events in the last day).

We also want to reject events too far in the future to avoid the current production issue.

Opposite to what we initially thought, this can already be deployed. **Additionally** we recommend to make a publish without recovery after deploying this version so that Pulse clock is effectively reset to "current time".

Performance improvement on Pulse Event Storage table access [🔗](#)

Feature

From now on Pulse Event Storage table is partitioned by timestamp in order to improve performance on read access. For now, it's implemented as monthly partition.

Fix Reference Data Service ansible roles tag [🔗](#)

Bug Fix

Before change:

```
- reference_data.415
- reference_data.configure
```

Copy

After change:

```
- reference-data.415
- reference-data.configure
```

Copy

Set JupyterLab memory limit [🔗](#)

Rework

Set JupyterLab memory limit to 8GB (`--ResourceUseDisplay.mem_limit=$((8*1024*1024*1024))`). We find this limit helpful to avoid unbounded memory usage by a single user. It is set to a fairly high value since, in our experience, it is uncommon that multiple users reach this maximum at the same time.

Bound the JupyterLab memory instead of the container in order to allow the software to self-manage its memory and give it a chance to optimize accordingly.

Ensure past events are considered by new plans real time and scheduled metrics [🔗](#)

Rework

By default a plan ignores past events that were already processed by the workflow. We've changed this default to account for our specific case, where we expect to deploy a plan soon and we want that plan to account for the data that is currently sent to Feedzai Pulse.

Fix duplicate feedback updates [🔗](#)

Bug Fix

Fixed an issue that caused multiple feedback updates to be created when marking an event as Fraud / Not Fraud in Case Manager.

Include Pulse Validation Tool [🔗](#)

Feature

The tool is able to validate that a Pulse Data Science Environment is configured as expected, by running common commands / jobs.

Run the tool using the `pulse-validation-tool.yml` playbook, present in `installer-assembly`. The playbook requires `Pulse Validation Tool`, present in [artifacts](#) page.

See Pulse Validation Tool Configurations for more information (later note: deprecated and removed in release `2.7.0`).

Ansible Roles Upgrade📦

Feature

Task "[921] create default users" has been moved to "[925] create default users", and task "[921] deploy security roles" has been added. Currently, this new task is not enabled.

Before:

- [921] create default users

After:

- [921] deploy security roles (currently skipped)
- [925] create default users