

sampling

Efficient sampling utilities through PySpark. This sub-package provides a stratified and a random sampling strategy, but can be easily extended to user defined strategies.

It supports smart optimized sampling that preserves the minimal history of entities, so historical profiles can be correctly computed while maximizing the sampling reduction power.

If you intend to implement your own sampling algorithm, that's fine: there's a `BaseSampler` that takes care of the integration with Pulse through the ds-api.